

OS Lab Assignment-1 (CS344)

Aditya Pandey (200123004)

Aman Bucha (200123006)

Sarthak Diwan (200123076)

Lab Group No: M20

Question-1

1. Created 3 syscalls, `thread_create`, `thread_join`, `thread_exit`.

Edited these files:

- a. syscall.c

```
extern int sys_draw(void);  
extern int sys_thread_create(void);  
extern int sys_thread_join(void);  
extern int sys_thread_exit(void);
```

```
[SYS_draw] sys_draw,  
[SYS_thread_create] sys_thread_create,  
[SYS_thread_join] sys_thread_join,  
[SYS_thread_exit] sys_thread_exit,
```

- b. syscall.h

```
#define SYS_draw 22  
#define SYS_thread_create 23  
#define SYS_thread_join 24  
#define SYS_thread_exit 25
```

- c. sysproc.c

```

int
sys_thread_create(void){
    int fcn;
    char* arg;
    char* stack;
    if(argint(0, &fcn) < 0) return -1;
    if(argint(1, (int*)&arg) < 0) return -1; //changed
    if(argint(2, (int*)&stack)<0) return -1; //changed

    return thread_create((void*)(void*))fcn, arg, stack);
}

```

```

int sys_thread_join(void)
{
    int ret_val = thread_join();
    return ret_val;
}

int sys_thread_exit(void){
    thread_exit();
    return 0;
}

```

d. defs.h

```

int      wait(void);
void     wakeup(void*);
void     yield(void);
int      thread_create(void(*)(void*), void*, void*);
int      thread_join(void);
void     thread_exit(void);

```

e. user.h

```

int sleep(int);
int uptime(void);
int thread_create(void(*)(void*), void*, void*);
int thread_join(void);
void thread_exit(void);

```

2. Implemented `thread_create` , `thread_join` and `thread_exit` in `proc.c`

Implementation:

int sys_thread_create(void)

It takes the 3 arguments (`fcn` , `arg` , `stack`) from the stack of the current process. Here `fcn` is the address of the function to execute. `arg` is the argument that will be passed to the function `fcn` . `stack` is the address of a memory region which will be used as **user stack** for the thread.

After taking the arguments it calls `thread_create` which is defined in `proc.c` .

int thread_create(void(*fcn)(void *), void * arg, void * stack)

It creates a new process by calling `allocproc` .

`allocproc` first acquires the lock for the `ptable` (the table which contains all the processes, defined as an array with 64 entries). It then searches the table for a process with state `UNUSED` , then changing the state to `EMBRYO` (meaning that stack is not yet allocated). It then sets up the kernel stack by allocating a free page using `kalloc` and then setting up space for `trapframe` , `trapret` and `context` .

In `trapframe` **userspace registers** are saved when the cpu changes from user mode to kernel mode. It will be loaded after `forkret` which is executed when the process starts executing.

`trapret` pops all the registers and save them.

`context` contains registers at the time of context switching.

Now in `thread_create` the page table of the new process(thread) is set as the same of the parent process. Now we set the `trapframe` of the new process.

`ESP` is the extended stack pointer which contains the address to the bottom of the stack. `stack` contains the address to the top of the memory allocated by `malloc` . Therefore, we set `esp` to `stack + 4096` . (4096 is the size of memory allocated).

`EIP` is the extended instruction pointer, which points to the next instruction to be executed. We set it to the address of the function `fcn` .

Now we pushed `arg` and `0xFFFFFFFF` (the return address) to the stack.

Lastly, the `cwd` (current working directory) and `name` are copied from the parent process and the process state is set to `RUNNABLE` (which means it can be scheduled by the scheduler). The ptable lock is released and `pid` is returned.

```
int
thread_create(void(*fcn)(void*), void* arg, void* stack)
{
    // how to pass argument in
    int i, pid;
    struct proc *np;
    struct proc *curproc = myproc();

    // Allocate process.
    if((np = allocproc()) == 0){
        return -1;
    }

    // Copy process state from proc.
    /*
    if((np->pgdir = copyvm(curproc->pgdir, curproc->sz)) == 0){
        kfree(np->kstack);
        np->kstack = 0;
    }
    */
}
```

```

    np->state = UNUSED;
    return -1;
}
*/

// this portion is particularly different
// alloting the page directory
np->pgdir = curproc->pgdir;
// alloting size of current process address space
np->sz = curproc->sz;
np->parent = curproc;
*np->tf = *curproc->tf;

// setting the instruction pointer to the addresss the passed function fcn
np->tf->eax = 0;
np->tf->eip = (uint)fcn;
// pointing the stack pointer to the end of the page allotted to it
np->tf->esp = (uint)stack+4096;

// similar to the fork function the child thread process also returns 0 to the parent thread/ process
np->tf->eax = 0;
// defining pop operation by traversing the memory
np->tf->esp -= 4;
*(uint*)(np->tf->esp) = (uint)(arg);

while(np->sz != curproc->sz)
{
    np->tf->esp -= 4;
    *(uint*)(np->tf->esp) = (uint)(arg);
    np->sz = curproc->sz;
    np->parent = curproc;
    *np->tf = *curproc->tf;
    if(!wait()) break;
}

// giuving a fake return address tp the function
np->tf->esp -= 4;
*(uint*)(np->tf->esp) = (uint)0xFFFFFFFF;

for(i = 0; i < NOFILE; i++)
    if(curproc->ofile[i])
        np->ofile[i] = filedup(curproc->ofile[i]);
np->cwd = idup(curproc->cwd);

safestrcpy(np->name, curproc->name, sizeof(curproc->name));

pid = np->pid;

acquire(&ptable.lock);

np->state = RUNNABLE;

release(&ptable.lock);

return pid;
}

```

int sys_thread_join(void)

It calls `thread_join` which is defined in `proc.c` .

int thread_join(void)

It first acquires the `ptable.lock` . Then it scans the process table to find the threads of the current process. For a process to be thread of the current process, its `parent` should be the current process **and** address of page directory of both the processes should be same. If such a process is not found then it returns `-1` . If a process is found with state `ZOMBIE` then it makes the process `UNUSED` and returns its pid. If the thread exists and not in `ZOMBIE` state, it waits for it to complete and get in the `ZOMBIE` state.

```
int
thread_join(void)
{
    struct proc *p;
    int havechild, pid;
    struct proc *curproc = myproc();

    acquire(&ptable.lock);
    while(1){
        // Scan through table looking for exited children.
        havechild = 0;
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->parent != curproc || p->pgdir != curproc->pgdir) // waiting for only threads and not forked processes to end
            {
                continue;
            }
            havechild = 1;
            if(p->state == ZOMBIE){
                // Found one.
                pid = p->pid;
                kfree(p->kstack);
                p->kstack = 0;
                //freevm(p->pgdir);
                p->parent = 0;
                p->pid = 0;
                p->state = UNUSED;
                p->name[0] = 0;
                p->killed = 0;
                release(&ptable.lock);
                return pid;
            }
        }

        // No point waiting if we don't have any children.
        if(!havechild || curproc->killed){
            release(&ptable.lock);
            return -1;
        }

        // Wait for children to exit. (See wakeup1 call in proc_exit.)
        sleep(curproc, &ptable.lock); //DOC: wait-sleep
    }
}
```

int sys_thread_exit(void)

It calls `thread_exit` which is defined in `proc.c` .

int thread_exit(void)

It gets the current process using `myproc` . It checks and panics if the current process is the `initproc` which is the first process which runs when `xv6` starts. Further it closes all the open files and resets the `curproc` . If parent does not exist, then control is passed to `initproc` . Finally the process state is set to `ZOMBIE` and scheduler is called.

```

void
thread_exit(void)
{
    struct proc *curproc = myproc();
    struct proc *p;
    int fd;

    if(curproc == initproc)
        panic("init exiting");

    // Close all open files.
    for(fd = 0; fd < NOFILE; fd++){
        if(curproc->ofile[fd]){
            fileclose(curproc->ofile[fd]);
            curproc->ofile[fd] = 0;
        }
    }

    begin_op();
    iput(curproc->cwd);
    end_op();
    curproc->cwd = 0;

    acquire(&ptable.lock);

    // Parent might be sleeping in wait().
    wakeup1(curproc->parent);

    // Pass abandoned children to init.
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->parent == curproc){
            p->parent = initproc;
            if(p->state == ZOMBIE)
                wakeup1(initproc);
        }
    }

    // Jump into the scheduler, never to return.
    curproc->state = ZOMBIE;
    sched();
    panic("zombie exit");
}

```

Result:

We can see that without using any mechanism like spinlocks/mutex the result (**3394**) is not the expected result (**6000**).

```

SeaBIOS (version 1.13.0-1ubuntu1.1)

ipXE (http://ipxe.org) 00:03.0 CA00 PCI2.10 PnP PMM+1FF8CA10+1FECCA10 CA00

ls          2 10 17756
mkdir       2 11 15368
rm          2 12 15348
sh          2 13 27992
stressfs    2 14 16260
usertests   2 15 67368
wc          2 16 17124
zombie      2 17 14940
Drawtest    2 18 15676
thread      2 19 19484
console     3 20 0
$ thread
SSttararitnign gd od_ow_owrokr:k :s: bs1:
b2
Done s:b2
Done s:b1
Threads finished: (5):6, (6):5, shared balance:3394
$ █

```

Question-2

For this question, we referred to the following lines in `spinlock.c` :

```

// Acquire the lock.
// Loops (spins) until the lock is acquired.
// Holding a lock for a long time may cause
// other CPUs to waste time spinning to acquire it.
void acquire(struct spinlock *lk)
{
    pushcli(); // disable interrupts to avoid deadlock.
    if (holding(lk))
        panic("acquire");

    // The xchg is atomic.
    while (xchg(&lk->locked, 1) != 0)
        ;

    // Tell the C compiler and the processor to not move loads or stores
    // past this point, to ensure that the critical section's memory
    // references happen after the lock is acquired.
    __sync_synchronize();

    // Record info about lock acquisition for debugging.
    lk->cpu = mycpu();
    getcallerpcs(&lk, lk->pcs);
}

// Release the lock.S
void release(struct spinlock *lk)

```

```

{
    if (!holding(lk))
        panic("release");

    lk->pcs[0] = 0;
    lk->cpu = 0;

    // Tell the C compiler and the processor to not move loads or stores
    // past this point, to ensure that all the stores in the critical
    // section are visible to other cores before the lock is released.
    // Both the C compiler and the hardware may re-order loads and
    // stores; __sync_synchronize() tells them both not to.
    __sync_synchronize();

    // Release the lock, equivalent to lk->locked = 0.
    // This code can't use a C assignment, since it might
    // not be atomic. A real OS would use C atomics here.
    asm volatile("movl $0, %0"
                  : "+m"(lk->locked)
                  :);

    popcli();
}

```

We tried to create a separate file for spinlocks and mutexes, wherein we had defined all the functions. However, we faced issues while including the c files, so we instead declared all the functions in `thread.c`

(a) Spinlocks

We defined the structure `thread_spinlock` as

```

struct thread_spinlock {
    uint locked;
};

```

We implemented the following functions:

void thread_spin_init(struct thread_spinlock *lk)

This function initializes the value of `lk->locked` to 0. This means the lock has not been acquired yet by any process.

void thread_spin_lock(struct thread_spinlock *lk)

This function waits for the lock to be freed. Implementation is done similar to the `acquire` function in `spinlock.c`. We first wait till the lock is acquired. Then, we call `__sync_synchronize`. This function tells CPU not to move loads or stores. This ensures critical section's memory references happen after the lock is acquired.

void thread_spin_unlock(struct thread_spinlock *lk)

We first call `__sync_synchronize`. Then we change the value of `lk->locked` to 0 (i.e. we release the lock). It was done using assembly language instruction to make it atomic.

Assembly Language Instruction: `asm volatile("movl $0, %0" : "+m" (lk->locked) :);not`

Result:

Now we are getting the expected result.

The final value of the total_balance matches the expected 6000, i.e., the sum of individual balances of each thread. This means that only one thread was allowed to enter the critical section at a time.

So implementation of spinlocks was done correctly.

```
console      3 20 0
$ thread
Ssttaarrttiinnngg ddoo__wworokr:k :s :sb:1
b2
Done s:b2
Done s:b1
Threads finished: (5):6, (6):5, shared balance:6000
$
```

(b) Mutexes

Implementation is done similar to spinlocks. The only difference is that we use the function `sleep` inside the `while` loop of `mutex_lock()` function.

We defined structure `mutex_lock` as

```
struct mutex_lock{
    uint locked;
};
```

We implemented the following functions:

void thread_mutex_init(struct thread_mutex *m)

This function initializes the value of `m->locked` to 0. This means the lock has not been acquired yet by any process.

void thread_mutex_lock(struct thread_mutex *m)

Similar to `thread_spin_lock` function, the difference being the use of `sleep` function inside the `while` loop. We use the sleep function because this releases CPU to another thread.

void thread_mutex_unlock(struct thread_mutex *m)

Function definition is exactly the same as that of `thread_spin_unlock` function, the only difference being the function parameter is `m` pointer not `lk` pointer.

Result:

Here too, we get the expected result.

The final value of the total_balance matches the expected 6000, i.e., the sum of individual balances of each thread. This means that only one thread was allowed to enter the critical section at a time.

So implementation of mutex locks was done correctly.

```
init: starting sh
$ thread
SStatratirntg idon_gw ordko: _sw:orkb1:
  s:b2
Done s:b1
Done s:b2
Threads finished: (4):4, (5):5, shared balance:6000
```